

Internet Congestion Control By Network Utility Maximization

Chee Wei Tan

Convex Optimization and its Applications to Computer Science

Outline

- Internet Congestion Control
- TCP Tahoe, TCP Reno and TCP Vegas
- Network Utility Maximization
- FAST TCP

TCP Congestion Control

- Window-based end-to-end flow control, where destination sends ACK for correctly received packets and source updates window size (which is proportional to allowed transmission rate)
- Several versions of TCP mitigate congestion by reducing window sizes. Sources update window sizes and links update (sometimes implicitly) congestion measures that are **feed back** to sources using the link

Optimization Model: TCP congestion control use **distributed algorithms** to solve an implicit **network utility maximization**, where source rates are **primal variables** updated at sources, and congestion measures are **dual variables** (shadow prices) updated at links

Why Congestion Control

Oct. 1986, Internet had its first congestion collapse (LBL to UC Berkeley)

- 400 yards, 3 hops, 32 kbps
- throughput dropped by a factor of 1000 to 40 bps

1988, Van Jacobson proposed TCP congestion control

- Window based with ACK mechanism
- End-to-end principle of Internet

Window-based Congestion Control

Limit number of packets in network to window size W

$$\text{Source rate allowed (bps)} = \frac{W \times \text{Message Size}}{\text{RTT}}$$

Too small W : **under-utilization** of link capacities

Too large W : link **congestion** occurs

Effects of congestion:

- Packet loss
- Retransmission and reduced throughput
- Congestion may continue after the overload

Basics of Congestion Control

- Goals: achieve high utilization without congestion or unfair sharing
- Receiver control (**awnd**): set by receiver to avoid overloading receiver buffer
- Network control (**cwnd**): set by sender to avoid overloading network
- $W = \min(\text{cwnd}, \text{awnd})$
- Congestion window cwnd usually the bottleneck

TCP Tahoe and Reno

- **Slow start**: start with window size 1 and doubles window size every RTT (increment window size by 1 per ACK)
- **Congestion avoidance**: pass a threshold, increase window by 1 every RTT (additive increase)
- Detecting a loss (Timeout or three duplicated ACK), retransmit loss packet, back to slow start, cut threshold by half
- **Fast recovery** in TCP Reno: between detecting loss and receiving ACK for retransmitted packet, temporarily increase window by 1 on receiving each duplicated ACK. After receiving ACK for retransmitted packet, cut window size by half and enters congestion avoidance directly

TCP Vegas

- Corrects **oscillatory** behavior of TCP Reno
- Obtains **queuing delay** by subtracting measured RTT from estimated propagation delay
- Sets rate proportional to ratio of propagation delay to queuing delay, proportionality constant between α and β
- Increase window if difference between queuing delay and propagation delay is smaller than α , decrease window if it is larger than β , remain constant otherwise

Queue Buffer Processes

At intermediate links:

- **FIFO** buffer process updates queuing delay as measure of congestion for Vegas and feeds back to sources
- **Drop tail** updates packet loss as measure of congestion for Reno and feeds back to sources
- **RED**: instead of dropping only at full buffer, drops packets with a probability that increases with (exponentially weighted) average queue length (example of **Active Queue Management**)

Analytic Models

Communication network with L links, each with fixed capacity c_l packets per second, shared by S sources, each using a set L_s of links

R : 0-1 routing matrix with $R_{ls} = 1$ iff $l \in L_s$

Deterministic flow model: $x_s(t)$ at each source s at discrete time t

Aggregate flow on link l :

$$y_l(t) = \sum_i R_{ls} x_s(t - \tau_{ls}^f)$$

where τ_{ls}^f is forward transmission delay

Each link updates congestion measure (shadow price) $p_l(t)$. Each

source has access to **aggregate price** along its route (**end-to-end**):

$$q_s(t) = \sum_l R_{ls} p_l(t - \tau_{ls}^b)$$

where τ_{ls}^b is backward delay in feedback path

Generic Source and Link Algorithms

Each source updates rate (z_s is a local state variable):

$$z_s(t+1) = F_s(z_s(t), q_s(t), x_s(t))$$

$$x_s(t+1) = G_s(z_s(t), q_s(t), x_s(t))$$

$$\text{Often } x_s(t+1) = G_s(q_s(t), x_s(t))$$

Each link updates congestion measure:

$$v_l(t+1) = H_l(y_l(t), v_l(t), p_l(t))$$

$$p_l(t+1) = K_l(y_l(t), v_l(t), p_l(t))$$

Notice access only to **local state** information (**distributed**)

Network Utility Maximization

Basic problem formulation:

$$\begin{array}{ll}\text{maximize} & \sum_s U_s(x_s) \\ \text{subject to} & Rx \preceq c \\ & x \succeq 0\end{array}$$

Objective: total utility (each U_s is smooth, increasing, concave)

Constraint: linear flow constraint

s index of sources and l index of links

Given routing matrix R_{ls} : 1 if flow from source s uses link l , 0 otherwise

x_s : source rate (**variables**), c_l : link capacity (**constants**)

Dual-based Distributed Algorithm

Use **optimization decomposition theory** to write **Lagrangian** $L(x, p)$:

$$\begin{aligned} &= \sum_s U_s(x_s) + \sum_l p_l \left(c_l - \sum_{s: l \in L(s)} x_s \right) \\ &= \sum_s \left[U_s(x_s) - \left(\sum_{l \in L(s)} p_l \right) x_s \right] + \sum_l c_l p_l \\ &= \sum_s L_s(x_s, q_s) + \sum_l c_l p_l \end{aligned}$$

Dual problem:

$$\begin{aligned} &\text{minimize} \quad g(p) = L(x^*(p), p) \\ &\text{subject to} \quad p \succeq 0 \end{aligned}$$

Dual-based Distributed Algorithm

Source algorithm:

$$x_s^*(q_s) = \operatorname{argmax} [U_s(x_s) - q_s x_s], \quad \forall s$$

- Selfish **net utility maximization** locally at source s

Link algorithm (gradient or subgradient based):

$$p_l(t+1) = \left[p_l(t) - \alpha(t) \left(c_l - \sum_{s:l \in L(s)} x_s^*(q_s(t)) \right) \right]^+, \quad \forall l$$

- Balancing supply and demand through **pricing**

Certain choices of step sizes $\alpha(t)$ (*e.g.*, $\alpha(t) = 1/t$) of **distributed algorithm** guarantee convergence to **globally optimal** (x^*, p^*)

Reverse Engineering TCP

KKT conditions:

- Primal and dual feasibility
- Lagrangian maximization (selfish net-utility maximization)
- Complementary slackness (generate the right prices to align selfish interest to social welfare maximization)

$p_l^* > 0$ indicates $y_l^* = c_l$ (link saturation) and

$y_l^* < c_l$ indicates $p_l^* = 0$ (buffer clearance)

Now specialize to **average model** of TCP Reno and TCP Vegas
Focus on **congestion avoidance** phase and **targeted equilibrium state**

TCP Reno: Source Algorithm

End-to-end marking probability q_s . Total delay D_s .

Window size w_s . Actual rate $\frac{w_s(t)}{D_s}$.

Net change to the window size:

$$x_s(t)(1 - q_s(t)) \cdot \frac{1}{w_s(t)} - x_s(t)q_s(t) \cdot \frac{1}{2} \cdot \frac{4w_s(t)}{3}.$$

Using $x_s = w_s/D_s$, we have

$$x_s(t+1) = \left[x_s(t) + \frac{1 - q_s(t)}{D_s^2} - \frac{2}{3}q_s(t)x_s^2(t) \right]^+.$$

TCP Reno: Arctan Utility

Arctan utility : $U_s(x_s) = \frac{\sqrt{3/2}}{D_s} \arctan \left(\sqrt{2/3} x_s D_s \right)$

Why?

At equilibrium ($t \rightarrow \infty$),

$$q_s = \frac{3}{2x_s^2 D_s^2 + 3}.$$

By optimality condition, need to check

$$U'_s(x_s) = q_s(x_s).$$

Get the utility function by integrating

TCP Vegas: Source Algorithm

Window size w_s

Propagation delay d_s . Expected rate $\frac{w_s(t)}{d_s}$

Queuing delay q_s and total delay D_s . Actual rate $\frac{w_s(t)}{D_s}$

$$w_s(t+1) = \begin{cases} w_s(t) + \frac{1}{D_s(t)} & \text{if } \frac{w_s(t)}{d_s} - \frac{w_s(t)}{D_s(t)} < \alpha_s \\ w_s(t) - \frac{1}{D_s(t)} & \text{if } \frac{w_s(t)}{d_s} - \frac{w_s(t)}{D_s(t)} > \alpha_s \\ w_s(t) & \text{else.} \end{cases}$$

Equilibrium round-trip time and window size satisfy:

$$\frac{w_s^*}{d_s} - \frac{w_s^*}{D_s^*} = \alpha_s$$

TCP Vegas: Log Utility Function

$$\text{Log utility} : U_s(x_s) = \alpha_s d_s \log x_s$$

Why: Complementary slackness condition is satisfied. Need to check

$$U'_s(x_s^*) = \frac{\alpha_s d_s}{x_s^*} = \sum_{l \in s} p_l^*$$

Let b_l^* be equilibrium backlog at link l . Window size equals bandwidth-delay product plus total backlog:

$$w_s^* - x_s^* d_s = \sum_{l \in s} \frac{x_s^*}{c_l} b_l^*$$

Using $x_s = w_s/D_s$, we have

$$\alpha_s = \frac{w_s^*}{d_s} - \frac{w_s^*}{D_s^*} = \frac{1}{d_s}(w_s^* - x_s^* d_s) = \frac{1}{d_s} \left(\sum_{l \in s} \frac{x_s^*}{c_l} b_l^* \right)$$

KKT condition satisfied if $p_l^* = \frac{b_l^*}{c_l}$ (dual variable is **queuing delay**)

TCP Vegas: Summary of Algorithm

Primal variable is source rate, updated by [source algorithm](#):

$$w_s(t+1) = [w_s(t) + v_s(t)]^+$$

$$v_s(t) = \frac{1}{d_s + q_s(t)} [\mathbf{1}(x_s(t)q_s(t) < \alpha_s d_s) - \mathbf{1}(x_s(t)q_s(t) > \alpha_s d_s)]$$

$$x_s(t) = \frac{w_s(t)}{d_s + q_s(t)}$$

Dual variable is queuing delay, updated by [link algorithm](#):

$$p_l(t+1) = \left[p_l(t) + \frac{1}{c_l} (y_l(t) - c_l) \right]^+$$

$$\text{Equilibrium: } x_s^* = \frac{\alpha_s d_s}{q_s^*}$$

TCP Reno and Vegas

TCP Reno (with Drop Tail or RED):

- Source utility: \arctan
- Link price: packet loss


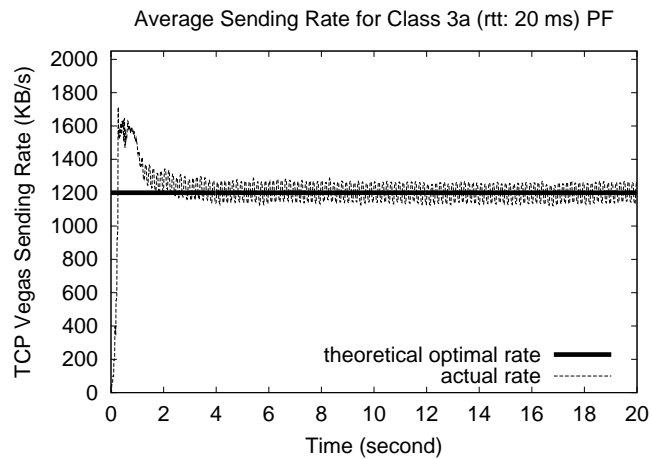
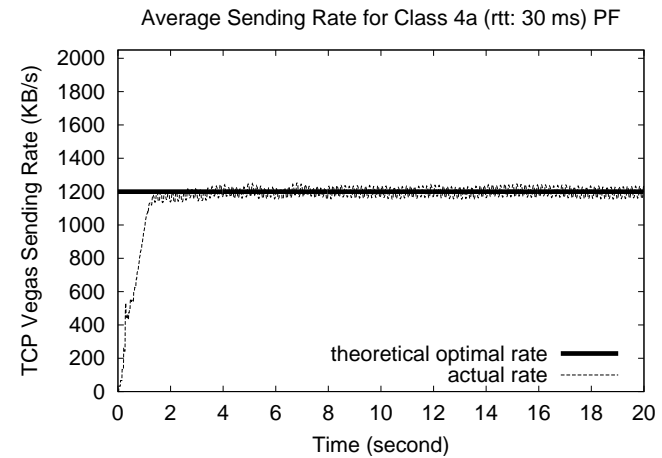
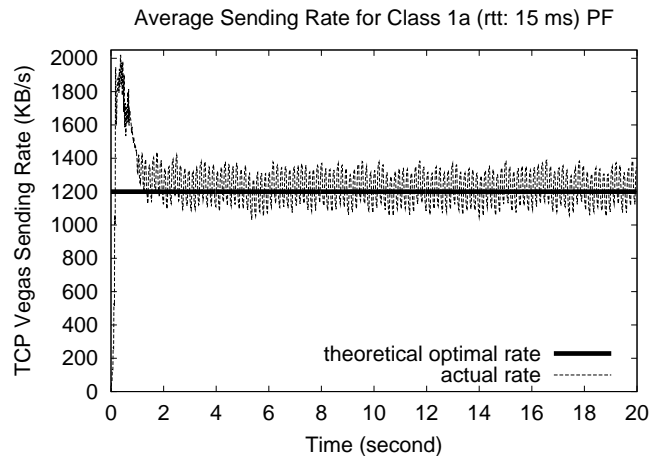
TCP Vegas (with FIFO)

- Source utility: weighted log
- Link price: queuing delay

Implications: Delay, Loss, Fairness

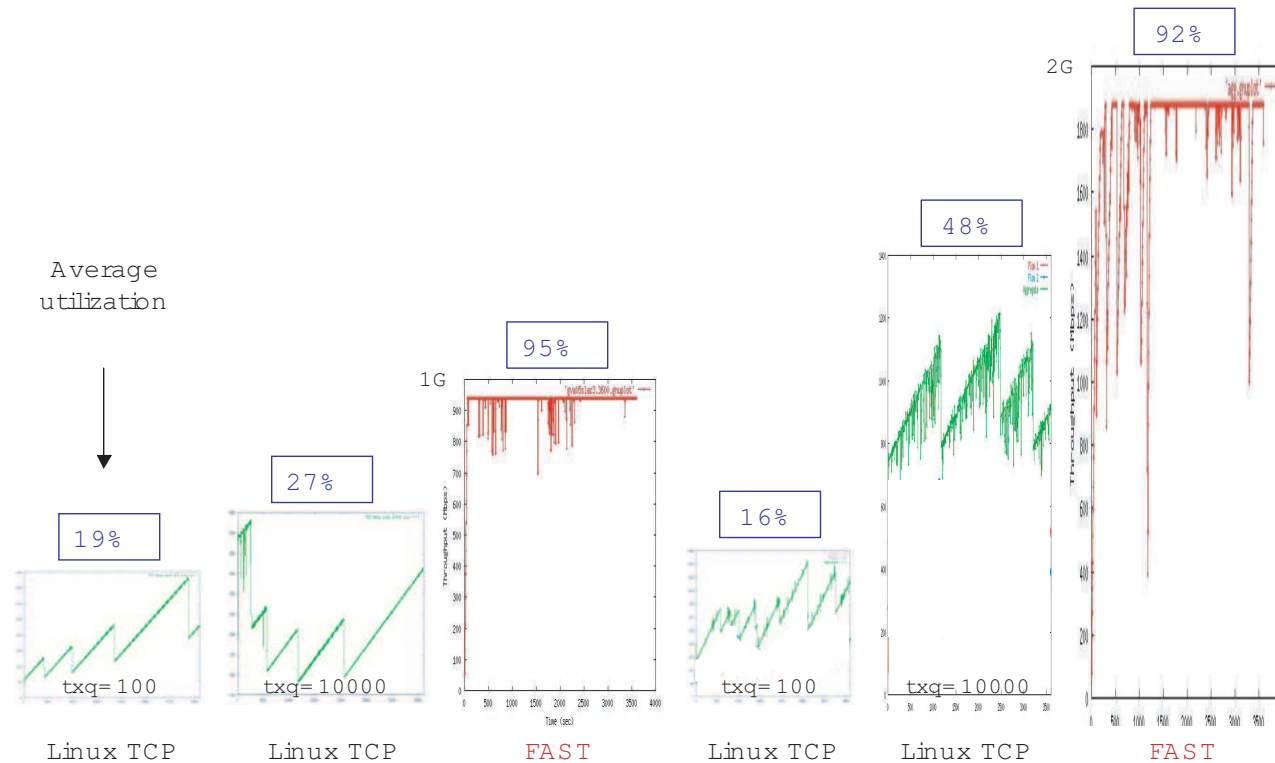
- TCP Reno: equilibrium loss probability is independent of link algorithms and buffer sizes. Increasing buffer sizes alone does not decrease equilibrium loss probability (buffer just fills up)
- TCP Reno: discriminates against connections with large propagation delays
- TCP Vegas: bandwidth-queuing delay product equals number of packets buffered in the network $x_s^* q_s^* = \alpha_s d_s$
- TCP Vegas: achieves proportional fairness
- TCP Vegas: gradient method for updating dual variable. Converges with the right scaling (γ small enough)

Numerical Example: Single Bottleneck



Successful Knowledge Transfer: FAST TCP

FAST TCP commercialized in 2006, increased bandwidth utilization efficiency in real networks from 20% to 90%, sold to Akamai in 2012



D. X. Wei, C. Jin, S. H. Low and S. Hegde, *FAST TCP: motivation, architecture, algorithms, performance*, IEEE/ACM Transactions on Networking, 2007

Summary

- Internet TCP congestion control as network optimizer: the use of *optimization criteria* and *optimization decomposition theory* to replace the prevalent “trial and error” approach to design
- Stability of TCP dual decomposition algorithm: In-class demo

Reading assignment:

- S. H. Low and D. E. Lapsley, Optimization flow control. I. Basic algorithm and convergence, IEEE/ACM Trans. Networking, 7(6), 1999
- S. Shakkottai and R. Srikant, Network Optimization and Control, Foundations and Trends in Networking, Now Publishers, 2007